

BIMM 143 Class 08: Mini Project

Xaler Lu (A17388454)

Background

In today's class, we will apply the methods of clustering techniques and PCA to make sense of a real-world breast cancer FNA biopsy data.

Setting up the data frame with `read.csv()` to read the csv file. Don't forget set the FIRST column as the name with `row.names=1` because we don't want the first column to be read as numeric.

```
fna.data <- "WisconsinCancer.csv"
wisc.df <- read.csv(fna.data, row.names = 1)
head(wisc.df)
```

	diagnosis	radius_mean	texture_mean	perimeter_mean	area_mean
842302	M	17.99	10.38	122.80	1001.0
842517	M	20.57	17.77	132.90	1326.0
84300903	M	19.69	21.25	130.00	1203.0
84348301	M	11.42	20.38	77.58	386.1
84358402	M	20.29	14.34	135.10	1297.0
843786	M	12.45	15.70	82.57	477.1
	smoothness_mean	compactness_mean	concavity_mean	concave.points_mean	
842302	0.11840	0.27760	0.3001	0.14710	
842517	0.08474	0.07864	0.0869	0.07017	
84300903	0.10960	0.15990	0.1974	0.12790	
84348301	0.14250	0.28390	0.2414	0.10520	
84358402	0.10030	0.13280	0.1980	0.10430	
843786	0.12780	0.17000	0.1578	0.08089	
	symmetry_mean	fractal_dimension_mean	radius_se	texture_se	perimeter_se
842302	0.2419	0.07871	1.0950	0.9053	8.589
842517	0.1812	0.05667	0.5435	0.7339	3.398
84300903	0.2069	0.05999	0.7456	0.7869	4.585
84348301	0.2597	0.09744	0.4956	1.1560	3.445

84358402	0.1809	0.05883	0.7572	0.7813	5.438
843786	0.2087	0.07613	0.3345	0.8902	2.217
area_se smoothness_se compactness_se concavity_se concave.points_se					
842302	153.40	0.006399	0.04904	0.05373	0.01587
842517	74.08	0.005225	0.01308	0.01860	0.01340
84300903	94.03	0.006150	0.04006	0.03832	0.02058
84348301	27.23	0.009110	0.07458	0.05661	0.01867
84358402	94.44	0.011490	0.02461	0.05688	0.01885
843786	27.19	0.007510	0.03345	0.03672	0.01137
symmetry_se fractal_dimension_se radius_worst texture_worst					
842302	0.03003	0.006193	25.38	17.33	
842517	0.01389	0.003532	24.99	23.41	
84300903	0.02250	0.004571	23.57	25.53	
84348301	0.05963	0.009208	14.91	26.50	
84358402	0.01756	0.005115	22.54	16.67	
843786	0.02165	0.005082	15.47	23.75	
perimeter_worst area_worst smoothness_worst compactness_worst					
842302	184.60	2019.0	0.1622	0.6656	
842517	158.80	1956.0	0.1238	0.1866	
84300903	152.50	1709.0	0.1444	0.4245	
84348301	98.87	567.7	0.2098	0.8663	
84358402	152.20	1575.0	0.1374	0.2050	
843786	103.40	741.6	0.1791	0.5249	
concavity_worst concave.points_worst symmetry_worst					
842302	0.7119	0.2654	0.4601		
842517	0.2416	0.1860	0.2750		
84300903	0.4504	0.2430	0.3613		
84348301	0.6869	0.2575	0.6638		
84358402	0.4000	0.1625	0.2364		
843786	0.5355	0.1741	0.3985		
fractal_dimension_worst					
842302	0.11890				
842517	0.08902				
84300903	0.08758				
84348301	0.17300				
84358402	0.07678				
843786	0.12440				

We want to omit the **diagnosis** column because that gives away the answer. To do this, use the code ‘df[,-1]’. Name this new data frame.

```
wisc.data <- wisc.df[,-1]
head(wisc.data)
```

	radius_mean	texture_mean	perimeter_mean	area_mean	smoothness_mean
842302	17.99	10.38	122.80	1001.0	0.11840
842517	20.57	17.77	132.90	1326.0	0.08474
84300903	19.69	21.25	130.00	1203.0	0.10960
84348301	11.42	20.38	77.58	386.1	0.14250
84358402	20.29	14.34	135.10	1297.0	0.10030
843786	12.45	15.70	82.57	477.1	0.12780
	compactness_mean	concavity_mean	concave.points_mean	symmetry_mean	
842302	0.27760	0.3001		0.14710	0.2419
842517	0.07864	0.0869		0.07017	0.1812
84300903	0.15990	0.1974		0.12790	0.2069
84348301	0.28390	0.2414		0.10520	0.2597
84358402	0.13280	0.1980		0.10430	0.1809
843786	0.17000	0.1578		0.08089	0.2087
	fractal_dimension_mean	radius_se	texture_se	perimeter_se	area_se
842302		0.07871	1.0950	0.9053	8.589 153.40
842517		0.05667	0.5435	0.7339	3.398 74.08
84300903		0.05999	0.7456	0.7869	4.585 94.03
84348301		0.09744	0.4956	1.1560	3.445 27.23
84358402		0.05883	0.7572	0.7813	5.438 94.44
843786		0.07613	0.3345	0.8902	2.217 27.19
	smoothness_se	compactness_se	concavity_se	concave.points_se	
842302	0.006399	0.04904	0.05373		0.01587
842517	0.005225	0.01308	0.01860		0.01340
84300903	0.006150	0.04006	0.03832		0.02058
84348301	0.009110	0.07458	0.05661		0.01867
84358402	0.011490	0.02461	0.05688		0.01885
843786	0.007510	0.03345	0.03672		0.01137
	symmetry_se	fractal_dimension_se	radius_worst	texture_worst	
842302	0.03003	0.006193	25.38	17.33	
842517	0.01389	0.003532	24.99	23.41	
84300903	0.02250	0.004571	23.57	25.53	
84348301	0.05963	0.009208	14.91	26.50	
84358402	0.01756	0.005115	22.54	16.67	
843786	0.02165	0.005082	15.47	23.75	
	perimeter_worst	area_worst	smoothness_worst	compactness_worst	
842302	184.60	2019.0	0.1622		0.6656
842517	158.80	1956.0	0.1238		0.1866
84300903	152.50	1709.0	0.1444		0.4245

```

84348301      98.87      567.7      0.2098      0.8663
84358402     152.20     1575.0      0.1374      0.2050
843786      103.40      741.6      0.1791      0.5249
               concavity_worst concave.points_worst symmetry_worst
842302          0.7119          0.2654          0.4601
842517          0.2416          0.1860          0.2750
84300903        0.4504          0.2430          0.3613
84348301        0.6869          0.2575          0.6638
84358402        0.4000          0.1625          0.2364
843786          0.5355          0.1741          0.3985
               fractal_dimension_worst
842302            0.11890
842517            0.08902
84300903          0.08758
84348301          0.17300
84358402          0.07678
843786            0.12440

```

Let's set up a new vector called `diagnosis` that has the diagnosis data.

```
diagnosis <- wisc.df$diagnosis
```

Exploring the Data Structure

Q1. How many observations are in this dataset?

There are 569 observations.

```
nrow(wisc.data)
```

```
[1] 569
```

Q2. How many of the observations have a malignant diagnosis?

The number of “M” malignant diagnosis is 212.

```
table(diagnosis)
```

```

diagnosis
  B    M
357 212

```

Q3. How many variables/features in the data are suffixed with `_mean`?

There are ten variables in the data with the “`_mean`” suffix. We use the `colname()` to list out all the column names, we then use `grep(pattern, data)` to count the number of the pattern “`_mean`”, and finally we use length to sum the number of patterns

```
colnames(wisc.data)
```

```
[1] "radius_mean"           "texture_mean"  
[3] "perimeter_mean"       "area_mean"  
[5] "smoothness_mean"       "compactness_mean"  
[7] "concavity_mean"        "concave.points_mean"  
[9] "symmetry_mean"         "fractal_dimension_mean"  
[11] "radius_se"             "texture_se"  
[13] "perimeter_se"          "area_se"  
[15] "smoothness_se"          "compactness_se"  
[17] "concavity_se"           "concave.points_se"  
[19] "symmetry_se"            "fractal_dimension_se"  
[21] "radius_worst"           "texture_worst"  
[23] "perimeter_worst"        "area_worst"  
[25] "smoothness_worst"       "compactness_worst"  
[27] "concavity_worst"        "concave.points_worst"  
[29] "symmetry_worst"         "fractal_dimension_worst"
```

```
length(grep("_mean", colnames(wisc.data)))
```

```
[1] 10
```

Principal Component Analysis

We want to make sure the data is scaled to the right unit of measurements and make sure the variables have significantly different variances because these could potentially violate our assumptions in statistical testing.

Check the means and SDs for ALL the columns with the function ‘

```
colMeans(wisc.data)
```

radius_mean	texture_mean	perimeter_mean
1.412729e+01	1.928965e+01	9.196903e+01
area_mean	smoothness_mean	compactness_mean
6.548891e+02	9.636028e-02	1.043410e-01
concavity_mean	concave.points_mean	symmetry_mean
8.879932e-02	4.891915e-02	1.811619e-01
fractal_dimension_mean	radius_se	texture_se
6.279761e-02	4.051721e-01	1.216853e+00
perimeter_se	area_se	smoothness_se
2.866059e+00	4.033708e+01	7.040979e-03
compactness_se	concavity_se	concave.points_se
2.547814e-02	3.189372e-02	1.179614e-02
symmetry_se	fractal_dimension_se	radius_worst
2.054230e-02	3.794904e-03	1.626919e+01
texture_worst	perimeter_worst	area_worst
2.567722e+01	1.072612e+02	8.805831e+02
smoothness_worst	compactness_worst	concavity_worst
1.323686e-01	2.542650e-01	2.721885e-01
concave.points_worst	symmetry_worst	fractal_dimension_worst
1.146062e-01	2.900756e-01	8.394582e-02

```
apply(wisc.data, 2, sd)
```

radius_mean	texture_mean	perimeter_mean
3.524049e+00	4.301036e+00	2.429898e+01
area_mean	smoothness_mean	compactness_mean
3.519141e+02	1.406413e-02	5.281276e-02
concavity_mean	concave.points_mean	symmetry_mean
7.971981e-02	3.880284e-02	2.741428e-02
fractal_dimension_mean	radius_se	texture_se
7.060363e-03	2.773127e-01	5.516484e-01
perimeter_se	area_se	smoothness_se
2.021855e+00	4.549101e+01	3.002518e-03
compactness_se	concavity_se	concave.points_se
1.790818e-02	3.018606e-02	6.170285e-03
symmetry_se	fractal_dimension_se	radius_worst
8.266372e-03	2.646071e-03	4.833242e+00
texture_worst	perimeter_worst	area_worst
6.146258e+00	3.360254e+01	5.693570e+02
smoothness_worst	compactness_worst	concavity_worst
2.283243e-02	1.573365e-01	2.086243e-01

concave.points_worst	symmetry_worst	fractal_dimension_worst
6.573234e-02	6.186747e-02	1.806127e-02

Execute PCA with `scale = T` to make sure the variance are not wildly different between each column

```
wisc.pr <- prcomp(wisc.data, scale = T)
summary(wisc.pr)
```

Importance of components:

	PC1	PC2	PC3	PC4	PC5	PC6	PC7
Standard deviation	3.6444	2.3857	1.67867	1.40735	1.28403	1.09880	0.82172
Proportion of Variance	0.4427	0.1897	0.09393	0.06602	0.05496	0.04025	0.02251
Cumulative Proportion	0.4427	0.6324	0.72636	0.79239	0.84734	0.88759	0.91010
	PC8	PC9	PC10	PC11	PC12	PC13	PC14
Standard deviation	0.69037	0.6457	0.59219	0.5421	0.51104	0.49128	0.39624
Proportion of Variance	0.01589	0.0139	0.01169	0.0098	0.00871	0.00805	0.00523
Cumulative Proportion	0.92598	0.9399	0.95157	0.9614	0.97007	0.97812	0.98335
	PC15	PC16	PC17	PC18	PC19	PC20	PC21
Standard deviation	0.30681	0.28260	0.24372	0.22939	0.22244	0.17652	0.1731
Proportion of Variance	0.00314	0.00266	0.00198	0.00175	0.00165	0.00104	0.0010
Cumulative Proportion	0.98649	0.98915	0.99113	0.99288	0.99453	0.99557	0.9966
	PC22	PC23	PC24	PC25	PC26	PC27	PC28
Standard deviation	0.16565	0.15602	0.1344	0.12442	0.09043	0.08307	0.03987
Proportion of Variance	0.00091	0.00081	0.0006	0.00052	0.00027	0.00023	0.00005
Cumulative Proportion	0.99749	0.99830	0.9989	0.99942	0.99969	0.99992	0.99997
	PC29	PC30					
Standard deviation	0.02736	0.01153					
Proportion of Variance	0.00002	0.00000					
Cumulative Proportion	1.00000	1.00000					

Q4. From your results, what proportion of the original variance is captured by the first principal component (PC1)?

PC1 captures 44.27% of the variance.

Q5. How many principal components (PCs) are required to describe at least 70% of the original variance in the data?

PC1 and PC2 is enough to explain at least 70% of the variance.

Q6. How many principal components (PCs) are required to describe at least 90% of the original variance in the data?

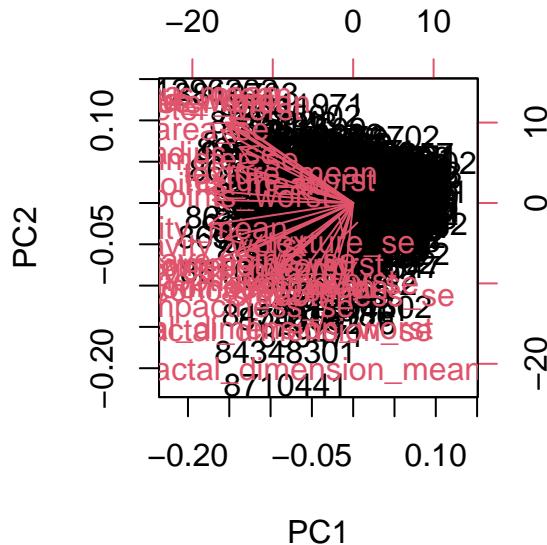
Up to PC7 is enough to explain at least 90% of the variance.

Interpreting PCA Results

Q7. What stands out to you about this plot? Is it easy or difficult to understand? Why?

This plot is a mess because I only see letters and numbers on the plot.

```
biplot(wisc.pr)
```



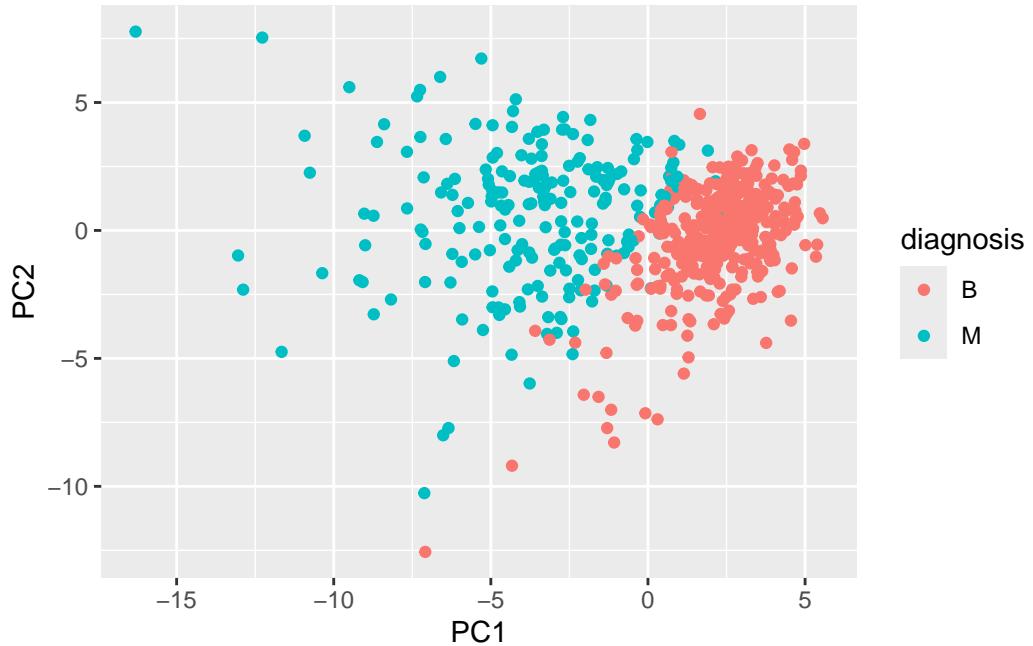
Let's use ggplot instead

```
library(ggplot2)
```

Warning: package 'ggplot2' was built under R version 4.3.3

PC1 VS PC2

```
ggplot(wisc.pr$x) +
  aes(PC1, PC2, col= diagnosis) +
  geom_point()
```

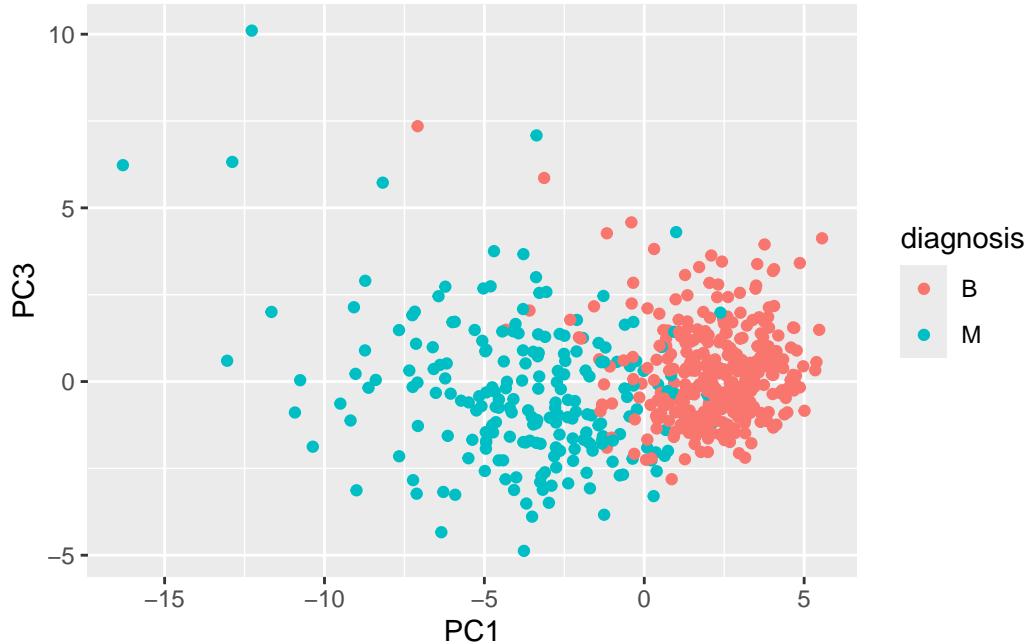


Q8. Generate a similar plot for principal components 1 and 3. What do you notice about these plots?

Nothing much changed, but I noticed that the outliers tend to shift up. Overall, the PC1 can discern the differences between benign and malignant. PC2 VS PC3 does not show a separation (not shown).

PC1 VS PC3

```
ggplot(wisc.pr$x) +
  aes(PC1, PC3, col= diagnosis) +
  geom_point()
```



Variance Explained

Use a scree plot to show the amount of variance each PC captures. Look for the “elbow” to show diminishing returns (usually with threshold 70% or 90%).

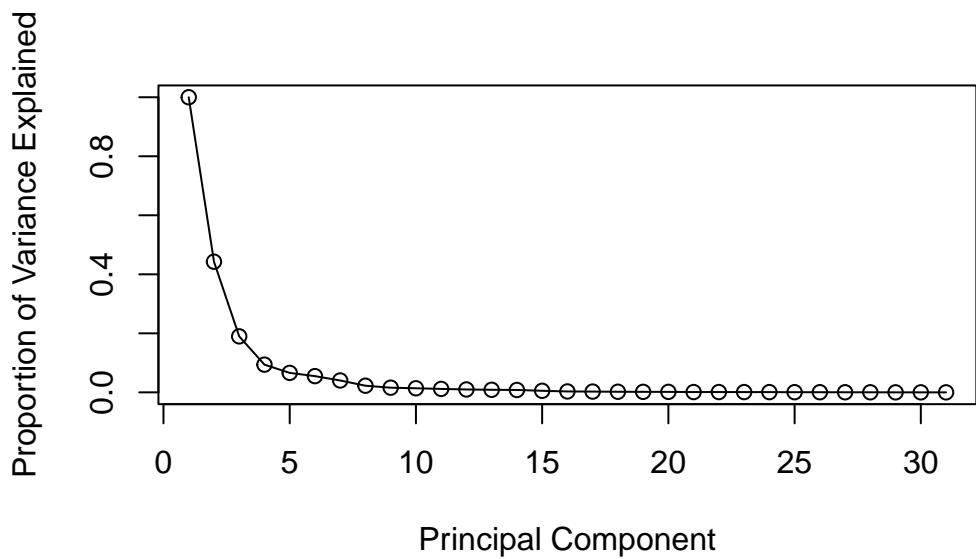
Recall variance is the square of SD.

```
pr.var <- wisc.pr$sdev^2
head(pr.var)
```

```
[1] 13.281608 5.691355 2.817949 1.980640 1.648731 1.207357
```

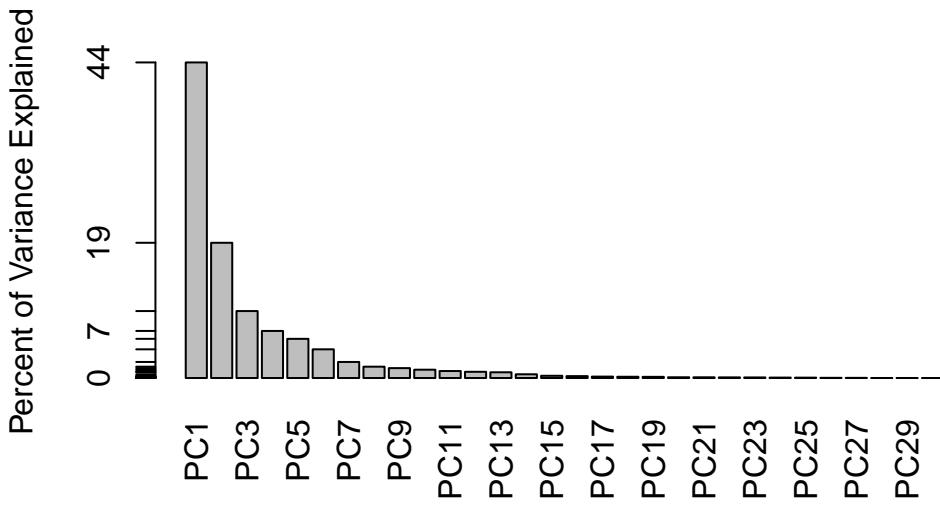
```
pve <- pr.var/sum(pr.var)

plot(c(1,pve), xlab = "Principal Component",
      ylab = "Proportion of Variance Explained",
      ylim = c(0, 1), type = "o")
```



An alternative scree plot

```
barplot(pve, ylab = "Percent of Variance Explained",
        names.arg=paste0("PC",1:length(pve)), las=2, axes = FALSE)
axis(2, at=pve, labels=round(pve,2)*100 )
```



Communicating PCA Results

Let's analyze the loading instead of the variance explained

Q9. For the first principal component, what is the component of the loading vector (i.e. `wisc.pr$rotation[,1]`) for the feature `concave.points_mean`? This tells us how much this original feature contributes to the first PC. Are there any features with larger contributions than this one?

The concave points mean for PC1 is -0.26. All the other measurements have a smaller negative (more positive) value than the concave point mean. Thus, nothing has more contribution than the concave points (in the negative direction).

```
wisc.pr$rotation["concave.points_mean", 1]
```

```
[1] -0.2608538
```

Hierarchical Clustering

Recall `hclust()` and its methods of clustering: `single`, `complete`, and `average`. Remember to scale `wisc.data` with `scale()`, and name this new data frame.

```
data.scaled <- scale(wisc.data)
```

Now, calculate the Euclidean distances with `dist()`

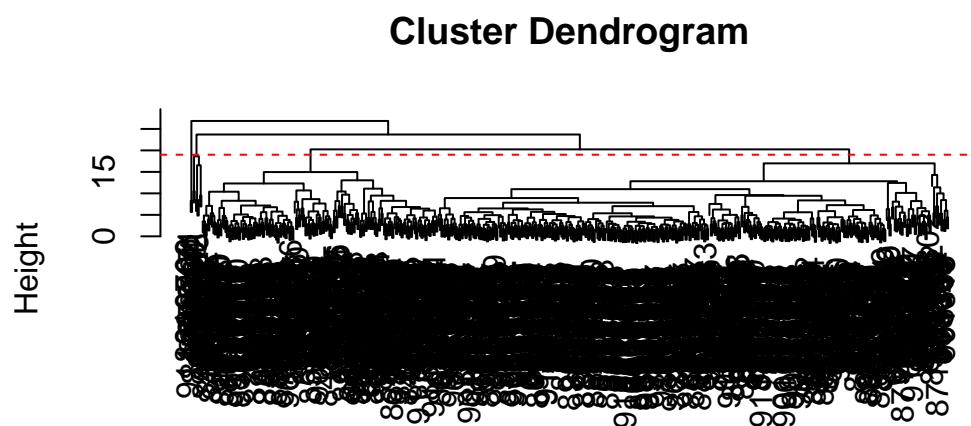
```
data.dist <- dist(data.scaled)
```

Use `hclust()` with the `complete` method

```
wisc.hclust <- hclust(data.dist, method = "complete")
```

Q10. Using the `plot()` and `abline()` functions, what is the height at which the clustering model has 4 clusters?

```
plot.hclust <- plot(wisc.hclust) +  
  abline(h = 19, col = "red", lty = 2)
```



Selecting the Number of Clusters

Cut the tree where the `abline` is using `cutree()`. Set the number of clusters `k` to 4

```
wisc.hclust.clusters.4 <- cutree(wisc.hclust, k = 4)
table(wisc.hclust.clusters.4, diagnosis)
```

diagnosis		
wisc.hclust.clusters.4	B	M
1	12	165
2	2	5
3	343	40
4	0	2

Clust into 6 groups

```
wisc.hclust.clusters.6 <- cutree(wisc.hclust, k = 6)
table(wisc.hclust.clusters.6, diagnosis)
```

diagnosis		
wisc.hclust.clusters.6	B	M
1	12	165
2	0	5
3	331	39
4	2	0
5	12	1
6	0	2

Q11. OPTIONAL: Can you find a better cluster vs diagnoses match by cutting into a different number of clusters between 2 and 6? How do you judge the quality of your result in each case?

Besides four clusters, none of the clusters between 2 to 6 better clusters the data.

Using Different Methods

Q12. Which method gives your favorite results for the same data.dist dataset?
Explain your reasoning.

Here is a function that specifies the number of clusters and the method for the wisc.hclust data.

```
wisc.hclust.method <- function(x, y) {
  wisc.hclust.a <- hclust(data.dist, method = y)
```

```
wisc.hclust.b <- cutree(wisc.hclust.a, k = x)
table(wisc.hclust.b, diagnosis)
}
```

Ward D2 is also my favorite method because it shows the separation with the least number of clusters.

```
wisc.hclust.method(2,"complete")
```

```
diagnosis
wisc.hclust.b   B   M
 1 357 210
 2   0   2
```

```
wisc.hclust.method(2,"single")
```

```
diagnosis
wisc.hclust.b   B   M
 1 357 210
 2   0   2
```

```
wisc.hclust.method(2,"average")
```

```
diagnosis
wisc.hclust.b   B   M
 1 357 209
 2   0   3
```

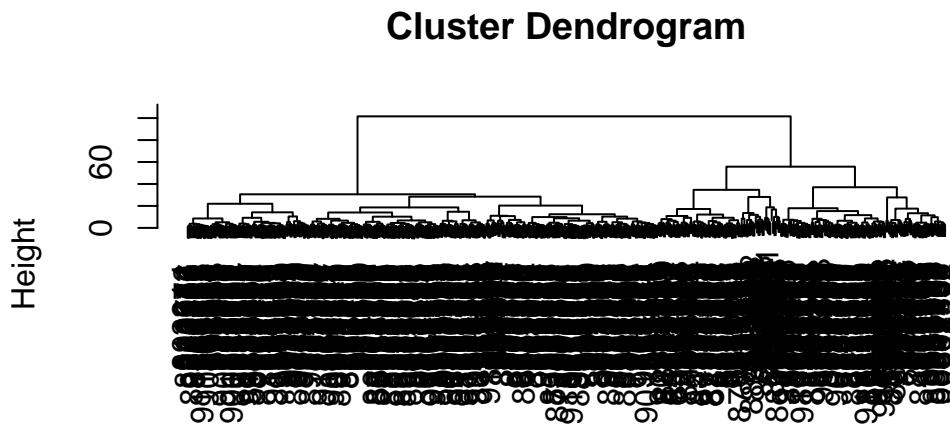
```
wisc.hclust.method(2,"ward.D2")
```

```
diagnosis
wisc.hclust.b   B   M
 1 20 164
 2 337 48
```

Combining Methods

Using the minimum number of PCs to describe 90% of the variability (PC1 through PC7), create a hierarchical clustering model with the method `ward.D2`. Instead of `hclust(dist(x))`, we use `hclust(dist(pr$x[,PCs]))`. In other words, we are combining both PCA and hierarchical clustering for this data analysis.

```
wisc.pr.hclust <- hclust(dist(wisc.pr$x[,1:7]), method = "ward.D2")
plot(wisc.pr.hclust)
```



```
dist(wisc.pr$x[, 1:7])
hclust (*, "ward.D2")
```

Let's cut this new tree. Cluster 1 mostly have M diagnosis and Cluster 2 has B diagnosis.

```
grps <- cutree(wisc.pr.hclust, k = 2)
table(grps)
```

```
grps
 1  2
216 353
```

```
table(grps, diagnosis)
```

```

diagnosis
grps   B    M
1     28 188
2    329  24

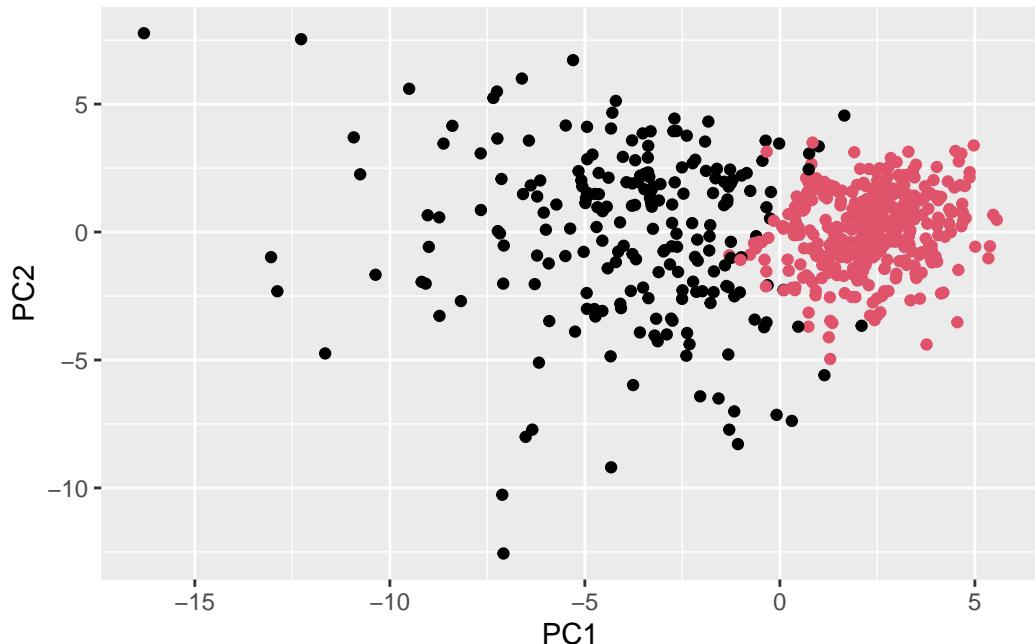
```

Let's make a ggplot.

```

ggplot(wisc.pr$x) +
  aes(PC1, PC2) +
  geom_point(col = grps)

```



Q13. How well does the newly created hclust model with two clusters separate out the two "M" and "B" diagnoses?

Without calculating the specificity or the selectivity, I do not see a noticeable difference between the hclust method and the PC + hclust method.

```

# v.name <- c("FP", "TP", "TN", "FN")
# sp.ward <- merge(wisc.hclust.method(2,"ward.D2"), v.name)
# sp.pr.ward <- merge(table(grps, diagnosis), v.name)
# sp.ward

```

Q14. How well do the hierarchical clustering models you created in the previous sections (i.e. without first doing PCA) do in terms of separating the diagnoses? Again, use the table() function to compare the output of each model (wisc.hclust.clusters and wisc.pr.hclust.clusters) with the vector containing the actual diagnoses.

```
wisc.hclust.method(2,"complete")
```

```
diagnosis
wisc.hclust.b   B   M
    1 357 210
    2   0   2
```

```
wisc.hclust.method(2,"single")
```

```
diagnosis
wisc.hclust.b   B   M
    1 357 210
    2   0   2
```

```
wisc.hclust.method(2,"average")
```

```
diagnosis
wisc.hclust.b   B   M
    1 357 209
    2   0   3
```

```
wisc.hclust.method(2,"ward.D2")
```

```
diagnosis
wisc.hclust.b   B   M
    1  20 164
    2 337  48
```

```
table(grps, diagnosis)
```

```

diagnosis
grps   B    M
1    28  188
2   329   24

```

Every hierarchical clustering method besides `ward.D2` did not cluster properly. The `complete` method also worked, but it needed four clusters to separate out the diagnosis.

Sensitivity and Specificity

Sensitivity = True Positive / (True Positive + False Negative)

Specificity = True Negative / (True Negative + False Positive)

Prediction

Use `predict()` to predice the old data `wisc.pr` and the new data `new`.

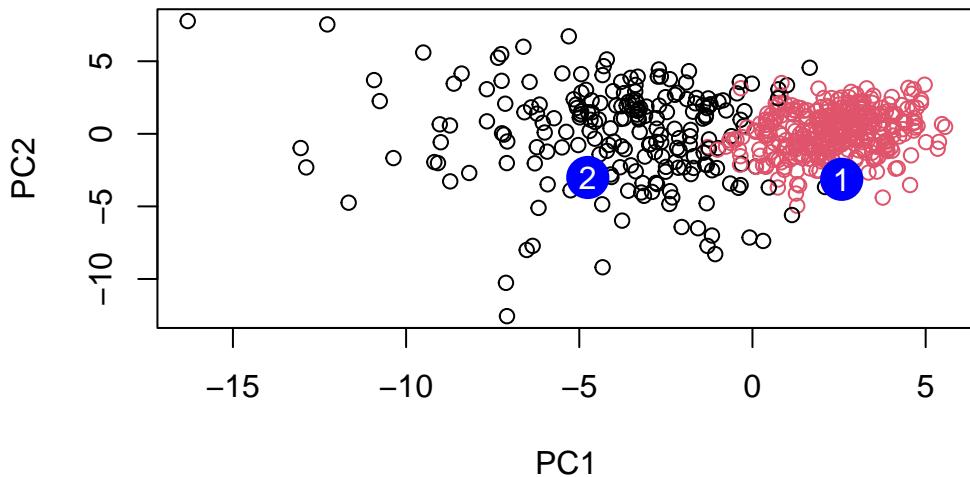
```

url <- "https://tinyurl.com/new-samples-CSV"
new <- read.csv(url)
npc <- predict(wisc.pr, newdata=new)
npc

```

	PC1	PC2	PC3	PC4	PC5	PC6	PC7
[1,]	2.576616	-3.135913	1.3990492	-0.7631950	2.781648	-0.8150185	-0.3959098
[2,]	-4.754928	-3.009033	-0.1660946	-0.6052952	-1.140698	-1.2189945	0.8193031
	PC8	PC9	PC10	PC11	PC12	PC13	PC14
[1,]	-0.2307350	0.1029569	-0.9272861	0.3411457	0.375921	0.1610764	1.187882
[2,]	-0.3307423	0.5281896	-0.4855301	0.7173233	-1.185917	0.5893856	0.303029
	PC15	PC16	PC17	PC18	PC19	PC20	
[1,]	0.3216974	-0.1743616	-0.07875393	-0.11207028	-0.08802955	-0.2495216	
[2,]	0.1299153	0.1448061	-0.40509706	0.06565549	0.25591230	-0.4289500	
	PC21	PC22	PC23	PC24	PC25	PC26	
[1,]	0.1228233	0.09358453	0.08347651	0.1223396	0.02124121	0.078884581	
[2,]	-0.1224776	0.01732146	0.06316631	-0.2338618	-0.20755948	-0.009833238	
	PC27	PC28	PC29	PC30			
[1,]	0.220199544	-0.02946023	-0.015620933	0.005269029			
[2,]	-0.001134152	0.09638361	0.002795349	-0.019015820			

```
plot(wisc.pr$x[,1:2], col=grps)
points(npc[,1], npc[,2], col="blue", pch=16, cex=3)
text(npc[,1], npc[,2], c(1,2), col="white")
```



Q16. Which of these new patients should we prioritize for follow up based on your results?

Patient one likely has a malignant diagnosis because they are in the malignant cluster according to PCA prediction.